

P0: Address Spaces and Resource Usage Monitoring

CSE 511, SPRING 2020

Due: February 4, 2020, 11:59:59pm

Praneet Soni and Bhuvan Urgaonkar

1 Goals

This programming assignment (P0) has three main goals. First, you will learn in concrete terms what the different segments of a virtual address space (or simply address space) are. Second, you will learn to use three Linux facilities - `/proc`, `strace`, and `getrusage` - to record/measure certain aspects of process resource usage. Finally, we hope that P0 will serve to refresh your knowledge of (or make you learn in case you lack it): (i) setting up and launching a compute environment on a public cloud platform (we will be working with AWS for this project) and using basic Linux shell and commands, (ii) compiling a C program (`gcc`, `make`), creating/linking against libraries, (iii) debugging (`gdb`), (iv) working with a code repository (`github`), (v) using Linux man pages (the `man` command), and (vi) plotting experimental data for easy visualization. All of these are good programming and experimental analysis practices that we would like you to follow throughout this course (and beyond it). A small guide to help you get started with AWS platform has been included at the end of this document.

2 Getting Started

After accepting the invitational link to the Github Classroom @CSE-511-SP20, you will be given access to your own private repository and another repository named "P0" containing all the files needed for completing P0. When you open this repository, you will find 5 folders, named `prog1`, ..., `prog5`. These folders contain the files described in Section 3. To access the Github Repository, use the following access link to join the classroom and a repository will be cloned for you : <https://classroom.github.com/a/wDvPG1iy>.

You are required to do the bulk of your work on a Linux (virtual) machine, deployed on a public cloud platform. The results in your report (which will be used for grading) should be based on your work on a cloud machine. This should be a good starting point, for you to familiarise yourself with the setup and configuration of a cloud environment that can facilitate your compute needs. However, to have relative consistency/uniformity in your measurements, you will need to use a **specific Amazon Machine Image** (packages containing the operating system and the requisite software binaries) for your

AWS instance. This would help us grade in a consistent manner and identify possible bugs/shortcomings. Please refer to the AWS starter guide at the end of the document carefully for further details. Adhere to the instructions there strictly, for the choice of your Virtual Machine.

Before answering the questions, make sure you disable Address Space Layout Randomization(ASLR) by running the following command:

```
echo 0 | sudo tee /proc/sys/kernel/randomize_va_space
```

However, be sure to disable ASLR everytime you start the AWS instance, as this will not persist across instance restarts.

3 Description of Tasks

1. **Stack, heap, and system calls:** The executable named `prog1` contains a function that is recursively called 10 times. This function has a local variable and a dynamically allocated variable. Upon each invocation, the function displays the addresses of the newly allocated variables on the console. After 10 invocations, the program waits for a key to be pressed on the keyboard before concluding. We would like you to observe the addresses displayed by `prog1` and answer the following.
 - (a) Which addresses are for the local variables and which ones are for the dynamically allocated variables? How were you able to deduce this? What are the directions in which the stack and the heap grow on your system?
 - (b) What is the size of the process stack when it is waiting for user input? (Hint: Use the contents of `/proc/PID/smmaps` that the `/proc` file system maintains for this process where we are denoting its process ID by `PID`. While the program waits for a user input, try running `ps -ef | grep prog1`. This will give you `PID`. You can then look at the `smmaps` entry for this process (`cat /proc/PID/smmaps`) to see a description of the current memory allocation to each segment of the process address space.
 - (c) What is the size of the process heap when it is waiting for user input?
 - (d) What are the address limits of the stack and the heap. (Hint: Use the `maps` entry within the `/proc` filesystem for this process. This will show all the starting and ending addresses assigned to each segment of virtual memory of a process.) Confirm the variables being allocated lie within these limits.
 - (e) Use the `strace` command to record the system calls invoked while `prog1` executes. For this, simply run `strace prog1` on the command line. Look at the man page of `strace` to learn more about it. Similarly, use man pages to learn basic information about each of these system calls. For each unique system call, write *in your own words* (just one sentence should do) what purpose this system call serves for this program.

2. **Debugging refresher:** The program `prog2.c` calls a recursive function which has a local and a dynamically allocated variable. Unlike the last time, however, this program will crash due to a bug we have introduced into it. Use the `Makefile` that we have provided to compile the program. Execute it. The program will exit with an error printed on the console. You are to compile the program and carry out the following tasks separately for each:
- (a) Observe and report the following for the 64 bit executables: (i) size of compiled code, (ii) size of code during run time, (iii) size of linked libraries.
 - (b) Use `gdb` to find the program statement that caused the error. See some tips on `gdb` in the Appendix if needed.
 - (c) Explain the cause of this error. Support your claim with address limits found from `/proc`.
 - (d) Using `gdb` back trace the stack. Examine individual frames in the stack to find each frame's size. Combine this with your knowledge (or estimate) of the sizes of other address space components to determine how many invocations of the recursive function should be possible on your system. How many invocations occur when you actually execute the program?
 - (e) What are the contents of a frame in general? Which of these are present in a frame corresponding to an invocation of the recursive function and what are their sizes?
3. **More debugging:** Consider the program `prog3.c`. It calls a recursive function which has a local and a dynamically allocated variable. Like the last time, this program will crash due to a bug that we have introduced in it. Use the provided `Makefile` to compile the program. Upon executing, you will see an error on the console before the program terminates. You are to carry out the following tasks:
- (a) Use `valgrind` to find the cause of the error including the program statement causing it. For this, simply run `valgrind ./prog3` on the command line. Validate this alleged cause with address space related information gleaned from `/proc`.
 - (b) How is this error different than the one for `prog2`?
4. **And some more:** The program `prog4.c` may seem to be error-free. But when executing under `valgrind`, you will see many errors. You are to perform the following tasks:
- (a) Describe the cause and nature of these errors. How would you fix them?
 - (b) Modify the program to use `getrusage` for measuring the following: (i) user CPU time used, (ii) system CPU time used - what is the difference between (i) and (ii)?, (iii) maximum resident set size - what is this?, (iii) signals received - who may have sent these?, (iv) voluntary context switches, (v) involuntary context switches - what is the difference between (iv) and (v)? Look at the sample code in the Appendix for an example on how to use `getrusage()`.

5. **Tracking resource usage: instrumenting the program vs. using external observations:** You are given executables for two programs (named `prog51` and `prog52`) that follow different regimes of dynamic memory allocations. Figures 1(a) and (b) depict fine-grained heap size evolution for these two programs, respectively. These were generated using a tool called Massif visualizer which instruments a program to record resource allocation information at run-time. You are to carry out the following tasks:

(a) Using appropriate system calls and scripting come up with your own solution for recording the evolution of the size of virtual memory allocated to a specified process. Your solution should record this for the process of interest once every second and output this into a text file. Here is an example of what such a file might look like:

```
7092
8134
12234
54874
345355
4347712
```

...

Use the plotting tool supplied by us (called `plot_script.py`) to create your own graphs of virtual memory allocation evolution for `prog51` and `prog52`. To run the plotting script:

```
$ python plot_script.py <path_to_your_file>
```

(b) Give a brief description of the difference between your graph and corresponding (included) graph. Why are you seeing this difference?

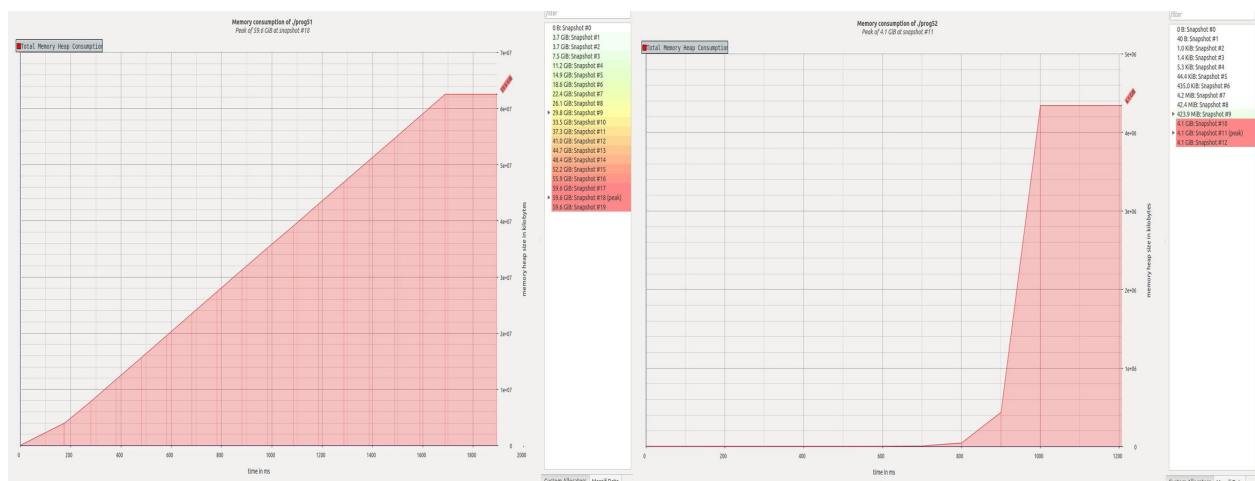


Figure 1: Heap allocations for `prog51` (left) and `prog52`.

4 Submission and Grading

You will submit all your source files, Makefiles, and READMEs (the last to convey something non-trivial we may need to know to use your code). You are to submit a report answering all the questions posed above into your github repo. Please **make sure** you commit all the required materials on your github repo before the deadline. Last commit at the time of deadline expiry will be used for the project evaluation and grading. P0 is worth 10 points (amounting to 3% of your overall grade). Each of the 5 tasks outlined above is worth 2 points. The TA will evaluate your submission based on (i) the quality and correctness of your report and code, and (ii) examining, compiling, and running code modified by you.

Appendix

We offer some useful hints here.

- **Quick notes on gdb:**
 1. To run a program `prog1` under `gdb`, simply execute
`$ gdb prog1`
 2. While running under `gdb`'s control, you can add breakpoints in the program to ease the debugging process. To add a breakpoint, type
`$ break <linenumber>`
 3. To run the code type
`$ r`
 4. To continue running the program after a breakpoint is hit, type
`$ c`
 5. To inspect the stack using `gdb`, type
`$ backtrace` or
`$ backtrace full` (to display contents of local variables)
 6. To get information about individual frames, type
`$ info frame <frame number>`
E.g., if you want to see information about frame 5 (assuming your program has made 6 recursive function calls, since frame number starts from 0), then the command would look like
`$ info frame 5`
 7. To get size of a frame, subtract frame addresses of two consecutive frames.
- To access virtual address space related information for a process with OS-assigned identifier `PID`, follow these steps:
 1. To find `PID` of, say, `prog2`, type,
`$ ps -ef | grep prog2` or,
`$ pgrep prog2`

2. Inspect the contents of the files `/proc/PID/maps` and `/proc/PID/smmaps` for a variety of useful information. A simple web search will offer details should you find something unclear (or ask us).
- Often a system's administrator will set an upper bound on the stack size. To find this limit:
`$ ulimit -s`
 Alternatively, you can use the following command to get both "soft" and "hard" limits set for a process:
`$ cat /proc/PID/limits`
 - To compile the code using 32/64 bit options, add the `-m<architecture>` flag to the compile command in the Makefile. E.g., to compile with the 32 bit option:
`$ gcc -g -m32 prog.c -o prog`
 - To find the size of an executable (including its code vs. data segments), consider using the `size` command. Look at its man pages.
 - To find the size of code during run time, type the following while the code is in execution:
`$ pmap PID | grep "total"`
 To see memory allocated to each section of the process, type
`$ pmap PID`
 - Sample code for using `getrusage()` :

```
#include <sys/time.h>
#include <sys/resource.h>
#include <unistd.h>
#include <stdio.h>

int main() {
    struct rusage usage;
    struct timeval start, end;
    int i, j, k = 0;

    getrusage(RUSAGE_SELF, &usage);
    start = usage.ru_stime;
    for (i = 0; i < 100000; i++) {
        for (j = 0; j < 100000; j++) {
            k += 1;
        }
    }
    getrusage(RUSAGE_SELF, &usage);
    end = usage.ru_stime;

    printf("Started at: %ld.%lds\n", start.tv_sec, start.tv_usec);
    printf("Ended at: %ld.%lds\n", end.tv_sec, end.tv_usec);
    return 0;
}
```

- **Starter Guide for Amazon EC2:** Amazon Elastic Compute Cloud (Amazon EC2) provides scalable computing capacity in the Amazon Web Services (AWS).

1. Sign up for an AWS account and you will be able to get started with Amazon EC2 for free. (750 Hours of free usage per month for 12 months)
2. Now to launch an instance, open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>
3. From the console dashboard, choose **Launch Instance**. Then you will be presented with a list of AMIs, that serve as templates for your instance. Select **Amazon Linux 2 AMI (HVM) 64-bit x86** , which will be marked as "Free tier eligible".
4. On the **Choose an Instance Type** page, you can select the hardware configuration of your instance. Select the t2.micro type, which is selected by default. Notice that this instance type is eligible for the free tier.
5. Choose Review and Launch to let the wizard complete the other configuration settings for you. On the Review Instance Launch page, under Security Groups, you'll see that the wizard created and selected a security group for you. We will use this security group for now. On the **Review Instance Launch** page, choose Launch.
6. You will then be prompted for a key pair. This key will be used for connecting to your instance and will authenticate your identity for access to your instance resources. Select **Create a new key pair**, enter a name for the key pair, and then choose Download Key Pair.
Note: This is the only chance for you to save the private key file, so be sure to download it. Save the private key file in a safe place. You'll need to provide the name of your key pair when you launch an instance and the corresponding private key each time you connect to the instance.
7. After you are done, choose Launch Instances. It takes a short time for an instance to launch. Choose View Instances to see the launch state. After the instance starts, it receives a public DNS name. We will use this public IP to connect to our instance.

- **Connect to your instance:** We will connect to the instance using an SSH client. You specify the private key (.pem) file, the user name for your AMI (which is `ec2-user` for our choice of AMI), and the public DNS name/IPv4 Public IP for your instance.

```
ssh -i /path/my-key-pair.pem ec2-user@IPv4_Public_IP
```

Note: To avoid incurring charges, Stop your instance when not in use. Terminate it once done with it, as that effectively deletes it and all the data will be lost. To do that, choose **Instances** and select the instance from the list. Now choose **Actions, Instance State** and then the required action.

- **Transferring files using SCP:** Transfer a file (For example, SampleFile.txt) to your instance using the instance's public DNS name/ IPv4 Public IP.

```
scp -i /path/my-key-pair.pem /path/SampleFile.txt  
ec2-user@ec2-198-51-100-1.compute-1.amazonaws.com:/path
```

Note: If **scp** is not already installed on the instance, do it using this :

```
sudo yum install -y openssh-clients
```

- **Additional information about EC2 instance:** Use the **yum** package manager to install all the requisite tools. A set of commonly needed developmental tools can be added to your instance by using

```
sudo yum groupinstall "Development Tools"
```

Use **yum** to install any additional tools needed, such as Valgrind. If needed by **valgrind**, **debuginfo** package can be downloaded using

```
yum install yum-utils
```

```
debuginfo-install glibc
```